



THE UNIVERSITY
ofADELAIDE



CRICOS PROVIDER 00123M

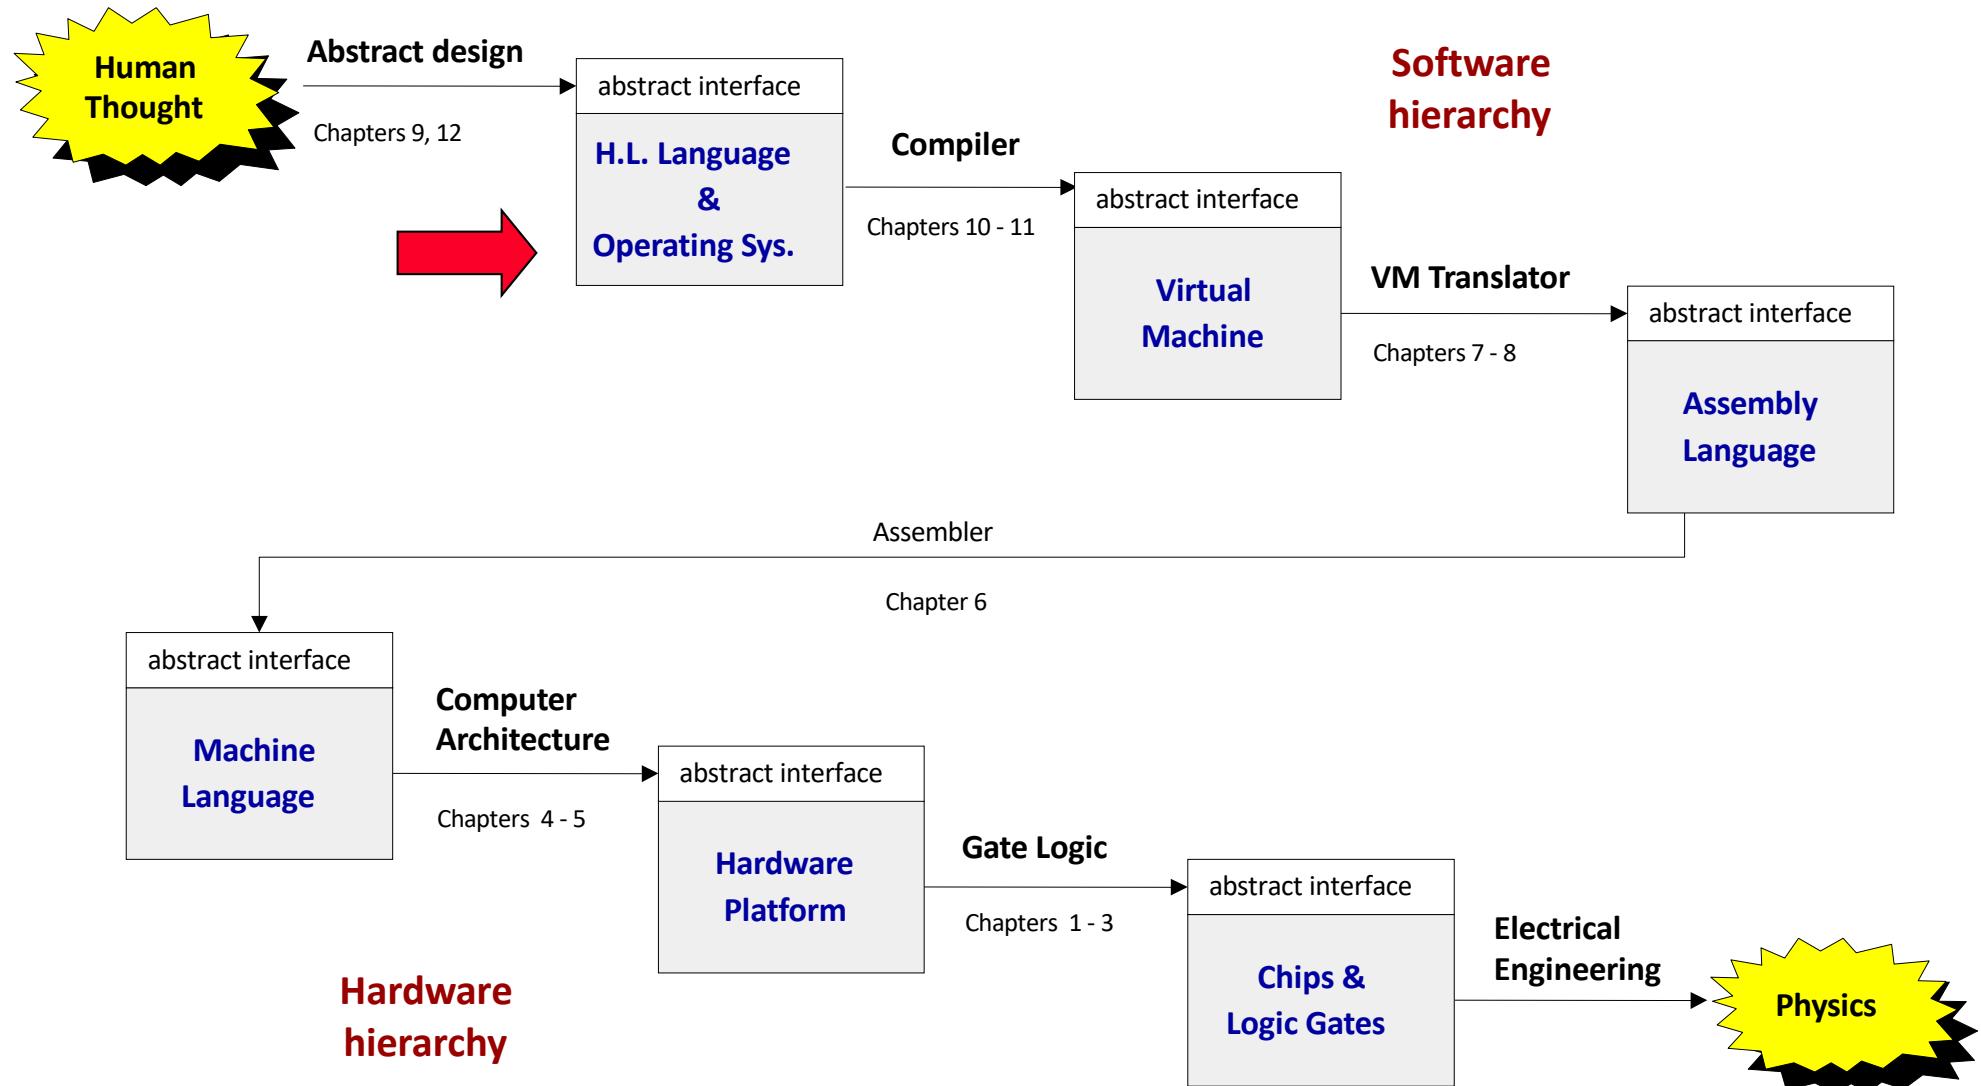
School of Computer Science

COMP SCI 2000 Computer Systems Lecture 22

adelaide.edu.au

seek LIGHT

Where we are at:



Lecture Plan

- The Operating System
 - Basic libraries
 - Single user system
- More efficient integer arithmetic
 - Multiply
 - Square root

Jack revisited

```
/** Computes the average of a sequence of integers. */
class Main
{
    function void main()
    {
        var Array a;
        var int length;
        var int i, sum;

        let length = Keyboard.readInt("How many numbers? ");
        let a = Array.new(length); // Constructs the array
        let i = 0;

        while (i < length)
        {
            let a[i] = Keyboard.readInt("Enter the next number: ");
            let sum = sum + a[i];
            let i = i + 1;
        }

        do Output.printString("The average is: ");
        do Output.printInt(sum / length);
        do Output.println();
        return;
    }
}
```

Typical OS functions

Language extensions / standard library

- Mathematical operations
(abs, sqrt, ...)
- Abstract data types
(String, Date, ...)
- Output functions
(printChar, printString ...)
- Input functions
(readChar, readLine ...)
- Graphics functions
(drawPixel, drawCircle, ...)
- And more ...

System-oriented services

- Memory management
(objects, arrays, ...)
- I/O device drivers
- Mass storage
- File system
- Multi-tasking
- UI management (shell / windows)
- Security
- Communications
- And more ...

The Jack OS

- **Math:** Provides basic mathematical operations;
- **String:** Implements the **String** type and string-related operations;
- **Array:** Implements the **Array** type and array-related operations;
- **Output:** Handles text output to the screen;
- **Screen:** Handles graphic output to the screen;
- **Keyboard:** Handles user input from the keyboard;
- **Memory:** Handles memory operations;
- **Sys:** Provides some execution-related services.

A typical OS:

- Is modular and scalable
- Empowers programmers (language extensions)
- Empowers users (file system, GUI, ...)
- Closes gaps between software and hardware
- Runs in “protected mode”
- Typically written in some high level language
- Typically grows gradually, assuming more and more functions
- Must be efficient.

Efficiency

We have to implement various operations on n -bit binary numbers ($n = 16, 32, 64, \dots$).

For example, consider *multiplication*

- Naïve algorithm: to multiply $x * y$: { for $i = 1 \dots y$ do $\text{sum} = \text{sum} + x$ }
 - Run-time is proportional to y
 - In a 64-bit system, y can be as large as 2^{64} .
 - Multiplications can take years to complete
- Algorithms that operate on n -bit inputs can be either:
 - Naïve: run-time is proportional to the value of the n -bit inputs
 - Good: run-time is proportional to n , the input's size.

Example I: multiplication

The “steps”

$$\begin{array}{r} 1 \ 0 \ 1 \ 1 \\ 1 \ 0 \ 1 \\ \hline 1 \ 0 \ 1 \ 1 \\ 0 \ 0 \ 0 \ 0 \\ 1 \ 0 \ 1 \ 1 \\ \hline 1 \ 1 \ 0 \ 1 \ 1 \end{array} = \begin{array}{r} 1 \ 1 \\ 5 \\ \hline \end{array}$$

The algorithm explained
(first 4 of 16 iteration)

$$\begin{array}{r} x: 0 \ 0 \ 0 \ 1 \ 0 \ 1 \ 1 \\ y: 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 1 \quad j\text{'th bit of } y \\ \hline 0 \ 0 \ 0 \ 1 \ 0 \ 1 \ 1 \\ 0 \ 0 \ 1 \ 0 \ 1 \ 1 \ 0 \\ 0 \ 1 \ 0 \ 1 \ 1 \ 0 \ 0 \\ 1 \ 0 \ 1 \ 1 \ 0 \ 0 \ 0 \\ \hline x \cdot y: 0 \ 1 \ 1 \ 0 \ 1 \ 1 \ 1 \end{array}$$

sum

multiply(x, y):

```
// Where  $x, y \geq 0$ 
sum = 0
shiftedX = x
for  $j = 0 \dots (n-1)$  do
    if ( $j$ -th bit of  $y$ ) = 1 then
        sum = sum + shiftedX
    shiftedX = shiftedX * 2
```

- Run-time: proportional to n
- Can be implemented in SW or HW
- Division: similar idea.

Example II: square root

The square root function has two convenient properties:

- It's inverse function is computed easily
- Monotonically increasing

Functions that have these two properties can be computed by binary search:

sqrt(x):

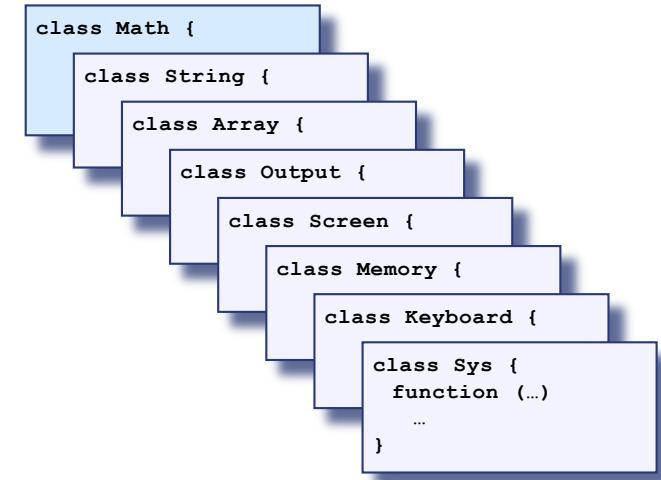
```
// Compute the integer part of  $y = \sqrt{x}$ . Strategy:  
// Find an integer  $y$  such that  $y^2 \leq x < (y+1)^2$  (for  $0 \leq x < 2^n$ )  
// By performing a binary search in the range  $0 \dots 2^{n/2} - 1$ .  
 $y = 0$   
for  $j = n/2 - 1 \dots 0$  do  
    if  $(y + 2^j)^2 \leq x$  then  $y = y + 2^j$   
return  $y$ 
```

Number of loop iterations is bounded by $n/2$, thus the run-time is $O(n)$.

Math operations

(in the Jack OS)

```
class Math {  
  
    function void init()  
  
    function int abs(int x)  
  
    ✓ function int multiply(int x, int y)  
  
    ✓ function int divide(int x, int y)  
  
    function int min(int x, int y)  
  
    function int max(int x, int y)  
  
    ✓ function int sqrt(int x)  
  
}
```



The remaining functions are simple to implement.

Perspective

- What we presented can be described as a:
 - mini OS
 - Standard library
- Many classical OS functions are missing
- No separation between user mode and OS mode
- Some algorithms (e.g. multiplication and division) are standard
- Other algorithms (e.g. line- and circle-drawing) can be accelerated with special hardware
- And, by the way, we've just finished building the computer.

Summary

- What have we done in this course?
 - We have seen how to build a computer system from Nand Gates up!
 - You have built
 - Basic gates
 - Sequential Logic
 - Assembler language programs
 - An Assembler
 - You have also studied
 - Boolean arithmetic circuits
 - Machine Language
 - CPU Design
 - Virtual Machines
 - VM-Translators
 - Compilers
 - OS's and High Level Languages

How did it all work?

- When you think about it a simple application written in Jack can translate to moderate number of VM instructions
- And these VM instructions translate into very many assembly language instructions.
- Once assembled these instructions ran on a CPU consisting of many hundreds of nand-gates
 - And memory consisting of over 100,000 gates.
- How did we handle this complexity?
 - Through modular design and
 - Well-designed interfaces
- By itself each component is not too complex
 - Wrap it in an interface and it becomes simpler still!
- These ideas of **interfaces** and **abstraction** are at the core of computer science.

Perspective

- Did we build a real machine?
 - Yes! The HDL can be mapped to hardware.
 - The software described would run on top of the machine as specified
- However
 - The Hack machine is designed for simplicity
 - Not for speed and scale
- Modern machines are more complex
 - Bigger word sizes, byte addressable, more complex and flexible instruction sets, more registers
 - Add to this a lot of complexity to speed up the machine
 - Caches, memory management hardware, instruction pipelines, branch prediction, out-of-order instruction issue, multiple cores...
- Most of this extra complexity is designed for extra performance
 - Very often – to make the common case fast
 - Another recurring theme in Computer Science!

Some Final notes

- CS is a science
- What is science?
- Reductionism
- Life science: From Aristo (millions of rules) to Darwin (3 rules) to Watson and Crick (1 rule)
- Computer science: We *knew* in advance that we could build a computer from almost nothing. In this course we actually did it.
- Key lessons:
 - Elegance
 - Clarity
 - Simplicity
 - Playfulness.



a	b	Out
0	0	1
0	1	1
1	0	1
1	1	0

