

STATS 2107

Statistical Modelling and Inference II

Practical 2: Manipulating Data

Sharon Lee, Matt Ryan

Semester 2 2022

In this practical, you start to look at how to

- manipulate the data,
- summarise data, and
- write functions.

Manipulating data

We are going to use the package `dplyr` to manipulate the `mpg` dataset. It is contained within the package `tidyverse` so let's load that:

```
library(tidyverse)
```

NOTE: This assumes `tidyverse` is installed, have a look at **Practical 1**.

Next, we load some data. I will use some built in data to make it easy to load. Load with

```
data("mpg")
mpg
```

```
## # A tibble: 234 x 11
##   manufacturer model      displ  year   cyl trans drv      cty   hwy fl      class
##   <chr>         <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi         a4          1.8  1999     4 auto~ f      18    29 p      comp~
## 2 audi         a4          1.8  1999     4 manu~ f      21    29 p      comp~
## 3 audi         a4          2    2008     4 manu~ f      20    31 p      comp~
## 4 audi         a4          2    2008     4 auto~ f      21    30 p      comp~
## 5 audi         a4          2.8  1999     6 auto~ f      16    26 p      comp~
## 6 audi         a4          2.8  1999     6 manu~ f      18    26 p      comp~
## 7 audi         a4          3.1  2008     6 auto~ f      18    27 p      comp~
## 8 audi         a4 quattro  1.8  1999     4 manu~ 4      18    26 p      comp~
## 9 audi         a4 quattro  1.8  1999     4 auto~ 4      16    25 p      comp~
## 10 audi        a4 quattro  2    2008     4 manu~ 4      20    28 p      comp~
## # ... with 224 more rows
```

Filter

How can I filter subjects given a criteria? For example, how could I get all cars manufactured by `audi`? In the `mpg` dataset, there is a column called `manufacturer` that has the information about each manufacturer. We can see the number of cars from each manufacturer with:

```
table(mpg$manufacturer)
```

```
##
##      audi  chevrolet    dodge    ford    honda  hyundai    jeep
##      18      19      37      25      9      14      8
## land rover    lincoln  mercury    nissan  pontiac    subaru    toyota
##      4        3        4        13      5        14      34
## volkswagen
##      27
```

So how do we get all the cars whose manufacturer is audi?

```
filter(mpg, manufacturer == "audi")
```

```
## # A tibble: 18 x 11
##   manufacturer model    displ  year  cyl trans drv   cty   hwy fl   class
##   <chr>          <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi          a4        1.8  1999   4 auto~ f    18    29 p    comp~
## 2 audi          a4        1.8  1999   4 manu~ f    21    29 p    comp~
## 3 audi          a4        2    2008   4 manu~ f    20    31 p    comp~
## 4 audi          a4        2    2008   4 auto~ f    21    30 p    comp~
## 5 audi          a4        2.8  1999   6 auto~ f    16    26 p    comp~
## 6 audi          a4        2.8  1999   6 manu~ f    18    26 p    comp~
## 7 audi          a4        3.1  2008   6 auto~ f    18    27 p    comp~
## 8 audi          a4 quattro 1.8  1999   4 manu~ 4    18    26 p    comp~
## 9 audi          a4 quattro 1.8  1999   4 auto~ 4    16    25 p    comp~
## 10 audi          a4 quattro 2    2008   4 manu~ 4    20    28 p    comp~
## 11 audi          a4 quattro 2    2008   4 auto~ 4    19    27 p    comp~
## 12 audi          a4 quattro 2.8  1999   6 auto~ 4    15    25 p    comp~
## 13 audi          a4 quattro 2.8  1999   6 manu~ 4    17    25 p    comp~
## 14 audi          a4 quattro 3.1  2008   6 auto~ 4    17    25 p    comp~
## 15 audi          a4 quattro 3.1  2008   6 manu~ 4    15    25 p    comp~
## 16 audi          a6 quattro 2.8  1999   6 auto~ 4    15    24 p    mids~
## 17 audi          a6 quattro 3.1  2008   6 auto~ 4    17    25 p    mids~
## 18 audi          a6 quattro 4.2  2008   8 auto~ 4    16    23 p    mids~
```

Let's break this down:

- `filter()` is the command that we want,
- `mpg` is the first argument and should be the name of the dataframe that we are going to filter, and
- `manufacturer == "audi"` is the next argument which gives the constraints in the filter command: note the use of `==` which is the logical equals.

More than one constraint - no problem:

```
filter(mpg, manufacturer == "audi", year <= 2000)
```

```
## # A tibble: 9 x 11
##   manufacturer model    displ  year  cyl trans drv   cty   hwy fl   class
##   <chr>          <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi          a4        1.8  1999   4 auto(~ f    18    29 p    comp~
## 2 audi          a4        1.8  1999   4 manua~ f    21    29 p    comp~
## 3 audi          a4        2.8  1999   6 auto(~ f    16    26 p    comp~
## 4 audi          a4        2.8  1999   6 manua~ f    18    26 p    comp~
## 5 audi          a4 quattro 1.8  1999   4 manua~ 4    18    26 p    comp~
## 6 audi          a4 quattro 1.8  1999   4 auto(~ 4    16    25 p    comp~
## 7 audi          a4 quattro 2.8  1999   6 auto(~ 4    15    25 p    comp~
```

```
## 8 audi          a4 quattro    2.8 1999      6 manua~ 4      17    25 p    comp~
## 9 audi          a6 quattro    2.8 1999      6 auto(~ 4      15    24 p    mids~
```

Quiz questions

1. What does the command given above return?
2. Filter for cars with front-wheel drive and 6 cylinders. How many cars satisfy these requirements?
3. What commands can be used to filter for all cars from `audi` or `dodge`?
4. Filter for all cars from `audi` or `dodge`. How many cars satisfy this condition?

Select

How do we choose a smaller set of columns? The command `select()` is your friend.

```
select(mpg, model, trans)
```

```
## # A tibble: 234 x 2
##   model      trans
##   <chr>     <chr>
## 1 a4        auto(15)
## 2 a4        manual(m5)
## 3 a4        manual(m6)
## 4 a4        auto(av)
## 5 a4        auto(15)
## 6 a4        manual(m5)
## 7 a4        auto(av)
## 8 a4 quattro manual(m5)
## 9 a4 quattro auto(15)
## 10 a4 quattro manual(m6)
## # ... with 224 more rows
```

If you want to select variables whose names contain a particular word, then use `contains()`:

```
select(mpg, contains("dis"))
```

```
## # A tibble: 234 x 1
##   displ
##   <dbl>
## 1  1.8
## 2  1.8
## 3    2
## 4    2
## 5  2.8
## 6  2.8
## 7  3.1
## 8  1.8
## 9  1.8
## 10    2
## # ... with 224 more rows
```

If you want to view a range of columns, use :

```
select(mpg, displ:cyl)
```

```
## # A tibble: 234 x 3
##   displ year  cyl
```

```
##      <dbl> <int> <int>
## 1  1.8  1999    4
## 2  1.8  1999    4
## 3  2    2008    4
## 4  2    2008    4
## 5  2.8  1999    6
## 6  2.8  1999    6
## 7  3.1  2008    6
## 8  1.8  1999    4
## 9  1.8  1999    4
## 10 2    2008    4
## # ... with 224 more rows
```

Quiz questions

5. What commands can be used to get the model, displ, and year.

Mutate

We can create new columns with the `mutate()` command. To illustrate, we are going to create a new column that is the city efficiency in km per litre.

```
mutate(mpg, cty_km_l = cty * 0.425144)
```

```
## # A tibble: 234 x 12
##   manufacturer model      displ  year   cyl trans drv      cty   hwy fl      class
##   <chr>          <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi          a4          1.8  1999    4 auto~ f      18    29 p      comp~
## 2 audi          a4          1.8  1999    4 manu~ f      21    29 p      comp~
## 3 audi          a4          2    2008    4 manu~ f      20    31 p      comp~
## 4 audi          a4          2    2008    4 auto~ f      21    30 p      comp~
## 5 audi          a4          2.8  1999    6 auto~ f      16    26 p      comp~
## 6 audi          a4          2.8  1999    6 manu~ f      18    26 p      comp~
## 7 audi          a4          3.1  2008    6 auto~ f      18    27 p      comp~
## 8 audi          a4 quattro  1.8  1999    4 manu~ 4      18    26 p      comp~
## 9 audi          a4 quattro  1.8  1999    4 auto~ 4      16    25 p      comp~
## 10 audi          a4 quattro  2    2008    4 manu~ 4      20    28 p      comp~
## # ... with 224 more rows, and 1 more variable: cty_km_l <dbl>
```

If you look at the end of the output, you can see the new column `cty_km_l`. So excited by this, let's look at mpg again:

```
mpg

## # A tibble: 234 x 11
##   manufacturer model      displ  year   cyl trans drv      cty   hwy fl      class
##   <chr>          <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi          a4          1.8  1999    4 auto~ f      18    29 p      comp~
## 2 audi          a4          1.8  1999    4 manu~ f      21    29 p      comp~
## 3 audi          a4          2    2008    4 manu~ f      20    31 p      comp~
## 4 audi          a4          2    2008    4 auto~ f      21    30 p      comp~
## 5 audi          a4          2.8  1999    6 auto~ f      16    26 p      comp~
## 6 audi          a4          2.8  1999    6 manu~ f      18    26 p      comp~
## 7 audi          a4          3.1  2008    6 auto~ f      18    27 p      comp~
## 8 audi          a4 quattro  1.8  1999    4 manu~ 4      18    26 p      comp~
## 9 audi          a4 quattro  1.8  1999    4 auto~ 4      16    25 p      comp~
## 10 audi          a4 quattro  2    2008    4 manu~ 4      20    28 p      comp~
```

```
## # ... with 224 more rows
```

ITS GONE!

The problem is that even though R created a new column, it did not save it. We need to tell R to do that:

```
mpg <- mutate(mpg, cty_km_l = cty * 0.425144)
```

Now it is saved.

Magrittr

What if you want to filter and select and mutate? Well you could do something like this:

```
mpg_audi <- filter(mpg, manufacturer == "audi")
mpg_trans_cty <- select(mpg_audi, trans:cty)
mpg_trans_cty_km <- mutate(mpg_trans_cty, cty_km_l = cty * 0.425144)
mpg_trans_cty_km
```

```
## # A tibble: 18 x 4
##   trans      drv    cty cty_km_l
##   <chr>    <chr> <int>   <dbl>
## 1 auto(l5)   f      18     7.65
## 2 manual(m5) f      21     8.93
## 3 manual(m6) f      20     8.50
## 4 auto(av)   f      21     8.93
## 5 auto(l5)   f      16     6.80
## 6 manual(m5) f      18     7.65
## 7 auto(av)   f      18     7.65
## 8 manual(m5) 4      18     7.65
## 9 auto(l5)   4      16     6.80
## 10 manual(m6) 4      20     8.50
## 11 auto(s6)   4      19     8.08
## 12 auto(l5)   4      15     6.38
## 13 manual(m5) 4      17     7.23
## 14 auto(s6)   4      17     7.23
## 15 manual(m6) 4      15     6.38
## 16 auto(l5)   4      15     6.38
## 17 auto(s6)   4      17     7.23
## 18 auto(s6)   4      16     6.80
```

This gets very annoying very quickly. Instead we can chain commands together using the command `%>%`. This command `%>%` is called a magrittr¹. Let's see an example:

```
mpg %>%
  filter(manufacturer == "audi") %>%
  select(trans:cty) %>%
  mutate(cty_km_l = cty * 0.425144)
```

```
## # A tibble: 18 x 4
##   trans      drv    cty cty_km_l
##   <chr>    <chr> <int>   <dbl>
## 1 auto(l5)   f      18     7.65
## 2 manual(m5) f      21     8.93
## 3 manual(m6) f      20     8.50
## 4 auto(av)   f      21     8.93
```

¹<https://cran.r-project.org/web/packages/magrittr/vignettes/magrittr.html>

```
## 5 auto(l5) f 16 6.80
## 6 manual(m5) f 18 7.65
## 7 auto(av) f 18 7.65
## 8 manual(m5) 4 18 7.65
## 9 auto(l5) 4 16 6.80
## 10 manual(m6) 4 20 8.50
## 11 auto(s6) 4 19 8.08
## 12 auto(l5) 4 15 6.38
## 13 manual(m5) 4 17 7.23
## 14 auto(s6) 4 17 7.23
## 15 manual(m6) 4 15 6.38
## 16 auto(l5) 4 15 6.38
## 17 auto(s6) 4 17 7.23
## 18 auto(s6) 4 16 6.80
```

Read `%>%` as `then`. So this reads as

- `mpg`: get the dataframe,
- `%>%`: then,
- `filter(manufacturer == "audi")`: filter for Audi,
- `%>%`: then,
- `select(trans:cty)`: select from `trans` to `cty`,
- `%>%`: then,
- `mutate(cty_km_l = cty * 0.425144)`: create a new column that contains the city efficiency in km per litre.

Challenge

Get all cars from `audi` that were made before or including 2000. Keep only the information for manufacturer, model, displ, year, cty. Filter these for all cars with a city efficiency of greater than 8 km/litre.

Quiz questions

6. How many cars passed the above challenge?

Group_by and summarise

To get summary statistics we can use `mean()` etc, but what if we wanted the mean city efficiency for each manufacturer? This can be done with a mix of `group_by` and then `summarise()`.

```
mpg %>%
  group_by(manufacturer) %>%
  summarise(mean = mean(cty, na.rm = TRUE))
```

```
## # A tibble: 15 x 2
##   manufacturer mean
##   <chr>         <dbl>
## 1 audi         17.6
## 2 chevrolet    15
## 3 dodge        13.1
## 4 ford         14
## 5 honda        24.4
## 6 hyundai      18.6
## 7 jeep         13.5
## 8 land rover   11.5
## 9 lincoln      11.3
```

```
## 10 mercury      13.2
## 11 nissan        18.1
## 12 pontiac      17
## 13 subaru       19.3
## 14 toyota       18.5
## 15 volkswagen   20.9
```

We can add more summary statistics:

```
mpg %>%
  group_by(manufacturer) %>%
  summarise(mean = mean(cty, na.rm = TRUE),
            n = n(),
            sd = sd(cty, na.rm = TRUE))
```

```
## # A tibble: 15 x 4
##   manufacturer mean      n    sd
##   <chr>         <dbl> <int> <dbl>
## 1 audi          17.6    18 1.97
## 2 chevrolet     15     19 2.92
## 3 dodge         13.1    37 2.49
## 4 ford          14     25 1.91
## 5 honda         24.4     9 1.94
## 6 hyundai       18.6    14 1.50
## 7 jeep          13.5     8 2.51
## 8 land rover    11.5     4 0.577
## 9 lincoln       11.3     3 0.577
## 10 mercury      13.2     4 0.5
## 11 nissan        18.1    13 3.43
## 12 pontiac      17      5 1
## 13 subaru       19.3    14 0.914
## 14 toyota       18.5    34 4.05
## 15 volkswagen   20.9    27 4.56
```

Functions

Often, you end up typing the same code again and again. Or at least copying and pasting a lot. Remember

Data analysis rule 3: Don't copy and paste - write a function.

So lets write a function to calculate the 95% confidence interval lower bound for the mean of some data.

First, type the following code into a script called `get_lwr_ci.R`.

```
# A function to take a vector of numbers and give a CI lower bound
#
# Uses the standard one for normal data with an unknown variance
# That is, this calculates the lower bound for a one sample test
# Using a t-distribution.
#
# arguments:
# x: a vector of numbers
# level: The strength of your confidence interval, i.e. 0.95 for 95%
#
# returns:
# vector with the lower point
get_lwr_ci <- function(x, level = 0.95){
```

```

# get mean
m <- mean(x, na.rm = TRUE)
# get SE
s <- sd(x, na.rm = TRUE)
n <- length(x)
se <- s / sqrt(n)
# get t cutoff point
a <- (1 - level) / 2
t <- qt(a, df = n-1, lower.tail = FALSE)
# lower point
lwr <- m - t * se
# return
return(lwr)
}

```

Save this in the same folder as your main script file.

Now in your main script file for this practical, you need to add the line:

```
source("get_lwr_ci.R")
```

NOTE: This assumes that the folder you have saved your main script file in is your working directory. Have a look at Practical 1 if you don't remember how to set this up.

When you run this line, it reads the file into the environment of R and if it is type correctly should give you a new function.

We can now test this with

```
get_lwr_ci(mpg$cty)
```

```
## [1] 16.31083
```

If you change your function, remember to rerun:

```
source("get_lwr_ci.R")
```

so that R loads the latest version into the environment.

Functions are of the basic form

```

name <- function(variables) {
  # Do stuff
  return(value)
}

```

The **name** is the name of the function, the variables are the arguments that will be passed into the function. Note that we can set default values like we did with `level = 0.95` in the `get_lwr_ci()` function. If we do not give a value for level, then 0.95 is used. Finally we return the values.

Challenge

1. Write the function to give the upper bound.
2. Get the 95% CI for the mean `cty` for each manufacturer.